

Master in HPC

Problem Sheet 4 - Quadtree

The aim of this problem is to construct a quadtree in two dimensions starting from a given set of N points (input file 'tree.dat') with coordinates $0 \leq x_i, y_i < L$ with $i = 1, \dots, N$

For these points write a program which computes the quadtree defined by the array `pointers_of_tree[0 : 3, root : root + NC]`, here NC is the number of cells of the tree (typically $NC \simeq N$) and $root = N + 1$ is the cell starting address.

Let us consider the square of side length L which encloses all of the points. The quadtree is constructed by calling a recursive function (`tree_sort`) which has as input the memory address of a generic cell and the list of particles which are enclosed within the cell itself. Each call to `tree_sort` increases the tree level `itercell_tree` by a unity. The tree construction begins by a call to `tree_sort` with `itercell_tree = 0`, and `root` as input cell together with the whole list of particles.

Before the first call the Morton keys of the particles are evaluated, setting `levorder = 10`.

A call to `tree_sort` finds first the subset of particles of the input list which fall within the boundaries of the four subcells. This is obtained by extracting for each particle the two bits of the Morton key which are distant $ndim \cdot itercell_tree$ positions from the top most significant bits. By construction, the integer corresponding to these bits represent the subcell index j . Particles of the input list with the same subcell index j are appended to a linked-list with header the subcell.

Let Np_{sub} be the number of points into each subcell. According to the value of Np_{sub} for the sub-quadrant j ; $j = 0, 1, 2, 3$, the integer value `inext` of the array `pointers_of_tree[j, root] = inext` takes the values :

$Np_{sub} > 1$ - more than one particle in the sub-quadrant, the sub-cell needs to be further examined : `inext = address new cell > root`.

$Np_{sub} = 1$ - there is one particle in the sub-quadrant, : `inext = i` address of the particle.

$Np_{sub} = 0$ - there are no particles in the sub-quadrant, there is no need to open further the sub-quadrant : `inext = 0`.

At each iteration the subcells with more than one particle are opened by

a new call to *tree_sort* which divides the subcell into four new subcells with size 1/2 of the parent cell. The procedure ends when there are no particles left to examine.

Write a program which reads as input file 'tree.dat' and for these points make a call to the function *tree_sort* which constructs the quadtree defined by *pointers_of_tree*[0 : 3, *root* : *root* + *NC*].

At the end write on a file *treecell.dat*, for all the cells of the quadtree, the four coordinates (bottom, top, left,right).

Write also on a file *treecell_part.dat* all the cell coordinates which contain one particle, that is the finall cells of the quadtree.

Finally write on a file *tree_ord.dat* the particles coordinates, ordered according to their Morton key-value.

Make a plots of the particles coordinates, together with the cells of the tree.

At the end of the pseudocodes you will find attached some useful references.

An efficient N-body algorithm for a fine-grain parallel computer. Use of Supercomputers in Stellar Dynamics; 1986, Springer-verlag

Hockney, R.W. & Eastwood, J.W., Computer Simulations uning Particles, Sect. 8.4, 1981, McGraw-Hill

Warren, M.S., A parallel hashed Oct-Tree N-body algorithm, Proceedings of the 1993 ACM/IEEE conference on Supercomputing, 1993

Here are given the corresponding pseudocodes.

Algorithm 1 Quadtree test

1: **procedure** TREE STRUCTURE ▷

Global:

Require: Int $Np = 4096$, $ncells = 2 * Np$, $nbodcell = Np + ncells$

Require: int $ndim = 2$, $nsubcell = 2^{ndim}$, $root = Np + 1$

Require: real $pos[2, nbodcell]$, $bottom[2, nbodcell]$, $cellsize[nbodcell]$

Require: int $pointers_of_tree[0 : nsubcell - 1, root : nbodcell]$

Require: int $iback[2, Np]$

Require: int $pm1[0 : nsubcell - 1, ndim]$

Require: int $subindex[Np]$, $bodlist[Np]$

Require: int $list_of_pointers[Np]$

Require: int $incells, N, itercell_tree, levorder$

Require: real side, rsize, rmin(ndim)

Local:

Require: int $jsub, ix, iy, ic, m, k, i, pcell$

Require: real $xcell[4]$, $ycell[4]$

Require: real $conv_to_int$ ▷ integer conversion

Require: long int $key_M[Np]$, $Morton_2D$ ▷ Morton key

Begin

2: Open *tree.dat* ▷ read data file

3: read side,N ▷ read box side and number of points

4: if $N > Np$ then

5: print N,Np

6: STOP

7: end if

8: for $i \leftarrow 1, N$ do

9: read $pos[1, i], pos[2, i]$ ▷ particles positions are stored in
 $pos[1 : 2, i]$

10: end for ▷ Now set the the sub-quadrant positions in pm1

11: $pm1[0, 1] := -1$

12: $pm1[0, 2] := -1$

13: $pm1[1, 1] := +1$

14: $pm1[1, 2] := -1$

15: $pm1[2, 1] := -1$

16: $pm1[2, 2] := +1$

17: $pm1[3, 1] := +1$

18: $pm1[3, 2] := +1$

```

19:      ▷ now compute the Morton keys - most significant bits to the right
20:      levorder := 10
21:      conv_to_int :=  $2^{levorder}/side$ 
21:      for i  $\leftarrow 1, N do
22:          ix = conv_to_int * pos[1, i]
23:          iy = conv_to_int * pos[2, i]
24:          key_M[i] = Morton_2D(ix, iy, levorder)
25:      end for
26:      rmin := 0                                ▷ set tree boundaries
27:      rsize := side
28:      incells := 1
29:      pointers_of_tree := 0
30:      cellsize := 0
31:      bottom := 0
32:      for k  $\leftarrow 1, ndim do                  ▷ set up position of root cell
33:          pos[k, root] := rmin[k] + rsize/2
34:          bottom[k, root] := rmin[k]
35:      end for
36:      cellsize[root] = side                  ▷ set up root cellsize
37:      for i  $\leftarrow 1, n do                      ▷ now initialize particle list
38:          bodlist[i] := i                      ▷ all particles need to be examined
39:      end for
40:      itercell_tree := 0                         ▷ set the level
41:      nbodlist := N                          ▷ actual value of the size of the particle list
42:      CALL tree_sort(nbodlist, bodlist, root)    ▷ call the tree function$$$ 
```

```

43: Open treecell.dat                                ▷ write cell file
44: print incells
45: for  $ic \leftarrow 1, incells$  do
46:      $p_{cell} := ic + root - 1$ 
47:      $x_{cell}[1] := bottom[1, p_{cell}]$ 
48:      $y_{cell}[1] := bottom[2, p_{cell}]$ 
49:      $x_{cell}[2] := bottom[1, p_{cell}] + cellsize[p_{cell}]$ 
50:      $y_{cell}[2] := bottom[2, p_{cell}]$ 
51:      $x_{cell}[3] := bottom[1, p_{cell}] + cellsize[p_{cell}]$ 
52:      $y_{cell}[3] := bottom[2, p_{cell}] + cellsize[p_{cell}]$ 
53:      $x_{cell}[4] := bottom[1, p_{cell}]$ 
54:      $y_{cell}[4] := bottom[2, p_{cell}] + cellsize[p_{cell}]$ 
55:     print  $ic, (x_{cell}[m], y_{cell}[m], m = 1, 4)$ 
56: end for

57: CALL sorti(key_M, subindex, N) ▷ this call returns in subindex the
   list of the particles ordered by the value of  $key\_M$ 
58: Open tree_ord.dat                               ▷ write particle file
59: print side
60: for  $m \leftarrow 1, N$  do
61:     i=subindex[m]
62:     print  $pos[1, i], pos[2, i], key\_M[i]$ 
63: end for
64: end procedure

```

```

1: procedure TREE_SORT(NBLIST,LISTBODIES,OLDCELL)
   Local:
Require: int  $j_{sub}$ ,  $j_u$ ,  $n_{subc}$ ,  $m$ ,  $i$ ,  $k$ ,  $k_{past}$ ,  $key\_of\_point$ 
Require: int  $p_{body}$ ,  $newcell$ ,  $oldcell$ ,  $levscan$ ,  $lpos$ 
Require: int  $listbodies[nblist]$ ,  $HOC[0 : 3]$ 
Require: int  $sublist[nblist]$ ,  $LLJ[nblist]$ 

2:    $itercell\_tree := itercell\_tree + 1$             $\triangleright$  monitor the level
3:   if  $itercell\_tree > levorder$  then
4:     print  $itercell\_tree$ ,  $levbit$        $\triangleright$  level higher than Morton order
5:     STOP
6:   end if
7:    $levscan := levorder - itercell\_tree$      $\triangleright$  bit positions for this level
8:    $HOC := 0$                                  $\triangleright$  reset the local linked-list
9:    $LLJ := 0$ 

10:  for  $k \leftarrow 1, nblist$  do           $\triangleright$  find the particle in each sub quadrant
11:     $i = listbodies[k]$                        $\triangleright$  which particle ?
12:     $key\_of\_point = key\_M[i]$              $\triangleright$  Morton key of the particle
13:     $lpos := ndim * levscan$   $\triangleright$  this is the starting address of the bits in
       the Morton key
14:     $ju := IBITS(key\_of\_point, lpos, 2)$    $\triangleright$  this is the sub-quadrant
15:     $k_{past} := HOC(ju)$                    $\triangleright$  append to the list
16:     $LLJ(k) := k_{past}$ 
17:     $HOC(ju) := k$ 
18:  end for
19:  for  $j_{sub} \leftarrow 0, n_{subcell} - 1$  do       $\triangleright$  now loop on the subcells
20:     $k := HOC(j_{sub})$                      $\triangleright$  first particle in the subcell
                                   $\triangleright$  subcell empty
21:    if  $k := 0$  then CYCLE
22:    end if
23:     $n_{subc} := 0$ 
24:    while  $k > 0$  do
25:       $n_{subc} := n_{subc} + 1$             $\triangleright$  copy the linked-list
26:       $i = listbodies[k]$ 
27:       $sublist[n_{subc}] = i$ 
28:       $k = llj[k]$ 
29:    end while                       $\triangleright$  end list

```

```

30:      if  $nsubc > 1$  then                                 $\triangleright$  this is a cell
31:           $incells := incells + 1$                    $\triangleright$  add the new cells
32:          if  $incells > ncells$  then
33:              print  $incells, ncells$ 
34:              STOP
35:          end if
36:           $newcell := incells + root - 1$             $\triangleright$  address new cell

37:           $pointers\_of\_tree[jsub, oldcell] := newcell$   $\triangleright \rightarrow$  pointer to the
   new cell
38:           $cellsize[newcell] := cellsize[oldcell]/2$ 
39:          for  $m \leftarrow 1, ndim$  do       $\triangleright$  new cell coordinates ( geometric
   center )
40:               $pos[m, newcell] := pos[m, oldcell] + pm1[jsub, m] * 0.5 *$ 
    $cellsize[newcell]$ 
41:               $bottom[m, newcell] := pos[m, newcell] - 0.5 *$ 
    $cellsize[newcell]$ 
42:          end for

43:           $iback[1, newcell] := jsub$      $\triangleright$  helpful to trace back the parent
   cell
44:           $iback[2, newcell] := oldcell$ 

45:          CALL  $tree\_sort(nsubc, sublist, newcell)$   $\triangleright$  go another level

46:      else if  $nsubc = 1$  then            $\triangleright$  one particle in the subcell
47:           $pbody := sublist[nsubc]$ 
48:           $pointers\_of\_tree[jsub, oldcell] := pbody$             $\triangleright$  because
    $pbody < root$  this indicates that  $pointers\_of\_tree[jsub, oldcell]$  points to a
   particle

49:           $iback[1, pbody] := jsub$ 
50:           $iback[2, pbody] := oldcell$ 
51:      end if
52:  end for                                 $\triangleright$  end subcells
53:   $itercell\_tree := itercell\_tree - 1$         $\triangleright$  back to the previous level
54: end procedure

```

An efficient N -body algorithm for a fine-grain parallel computer

Joshua E. Barnes

Institute for Advanced Study
Princeton, NJ 08540

This paper describes an implementation of a hierarchical N -body force-calculation algorithm on a giant electronic brain with $O(N)$ processors. For large- N problems this approach appears to offer an optimal scaling of real computational cost with N . Other problems involving long-range interactions in highly inhomogeneous systems may also be solved with these techniques.

1. Introduction

Numerical simulations of complicated N -body systems require codes with significantly greater capacity than provided by existing schemes. For example, to study the collision of two galaxies with realistic self-gravitating disks, we will probably need to integrate the equations of motion of $\sim 2.5 \times 10^5$ particles. Problems of this kind require algorithms which can handle a wide range of physical scales without any special assumptions about the symmetry of the system. Conventional algorithms have difficulty meeting these requirements: direct-summation codes are far too slow, FFT codes can handle only a limited range of scales, and multipole-expansion codes make restrictive assumptions about the geometry. This report describes an N -body algorithm for a parallel computer combining the generality of direct-summation codes with the favorable scaling properties of FFT and multipole-expansion codes.

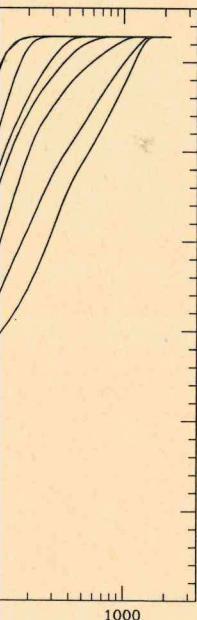
2. Force calculation on a tree

The most time-consuming part of a direct-summation N -body code is always the force calculation. To calculate the total forces acting on N particles requires $O(N^2)$ operations, since each particle interacts with $N - 1$ others. Recently, Barnes & Hut (1986) have described an algorithm which approximates the force on a particle by computing only $O(\log N)$ interactions. A complete force calculation thus takes $O(N \log N)$ operations, a significant improvement for large N .

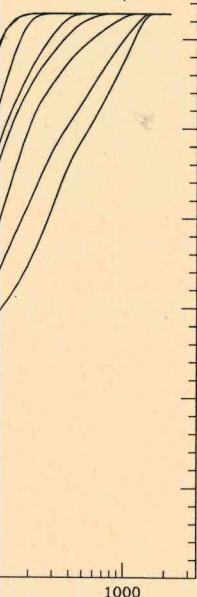
The new algorithm employs a strategy similar to that used by the finite human brain to deal with an apparently infinite universe. For example, it is clearly important to take the wishes and behavior of other human beings into account in planning one's actions. Yet anyone who attempted to consider $\sim 5 \times 10^9$ individuals would face an impossible computational task. The solution is to adopt a lumped description, treating "nearby" people as individuals, but grouping more distant peoples into ever larger units. For this strategy to work, each individual must adopt a different description, and these descriptions must change with time.

These requirements dictate many of the features of the force-calculation algorithm. The description used for particle y cannot be re-used for particle x . Furthermore, when calculating the force on x , it is not permissible to scan the entire set of particles and decide which to lump together, since that would take time of $O(N)$ and so destroy the fundamental advantage of the algorithm. Instead, a tree-structure is used to represent a set of possible lumped descriptions of the N -body system. This tree is generated from the particle positions at the beginning of each time-step. The leaves of the tree are the particles; interior nodes represent groups of particles which are physically near each other and may be lumped together. The force on a particle is calculated by recursively descending the tree, starting at the "root" node, which represents the group of all particles. The recursion terminates whenever it reaches a leaf (*i.e.*, particle), or

Example 1



Example 2



an internal node representing a group of particles sufficiently distant that the error incurred in lumping them together can be accepted.

An implementation of this algorithm on a conventional serial processor is discussed by Barnes & Hut. For a system of $N = 4096$ particles, forces can be calculated to $\sim 1\%$ accuracy in a tenth of the time required by a direct-summation algorithm. CPU time scales very closely to the expected $N \log N$ for $N = 512$ to $32,768$. To obtain the full advantage of a tree-code, however, we want to consider problems involving much larger values of N . For $N \simeq 10^5$ to 10^6 , two "economic" considerations arise. First, conventional serial processors do not have the floating-point throughput necessary to run interesting calculations in reasonable time. Typical vector or pipelined processors do not lend themselves to tree-codes, since numeric operations are interwoven with logical decisions. Second, as conventional computers are made larger and larger, less and less of the hardware is devoted to the CPU, and more and more is used for memory. In the large- N limit, the total cost of an experiment scales like $N^2 \log N$, since the expense of "renting" the hardware also scales with N .

3. Fine-grain parallel architectures

Why should one rethink the basic choice of computer architecture, and what are the advantages of a fine-grain approach? The answer to the first question is partly historical. The architecture adopted by von Neumann and his collaborators reflected the computational priorities and economic factors of the day. A simple and general structure which could sustain a reasonable rate of computation was favored over a design which maximized the rate of basic logical operations. Reasonably fast processing demanded vacuum tubes, which were an expensive and therefore centralized resource; memory, on the other hand, was relatively slow and inexpensive. These factors dictated a two-part structure, with a fast CPU accessing inert memory banks through a single high-speed channel. Today, the economics of computer hardware have changed, and the central processing unit and its associated memory are both made out of silicon. It has become apparent that the von Neumann architecture makes very uneven use of the processing power inherent in raw silicon. Within the CPU, a large number of switching elements change state with every clock cycle, while within the memory, only of order unity do. Large von Neumann computers typically exploit the raw potential of the hardware with an efficiency of $\sim 3\%$ or less.

Fine-grain parallel architectures offer a way out of this "von Neumann bottleneck". In machines of this kind, storage and computational units are combined in processing cells which are put together in a homogeneous pattern, and a large fraction of the hardware can be active each clock cycle. As Hillis (1985) observes, the construction of such a computer is rather straightforward once the basic design has been worked out. The real problem is to come up with a architecture with a degree of flexibility comparable to that of the von Neumann architecture. Hillis identifies two general requirements for such a design. First, to make the job of programming tractable, it must be possible to set up a one-to-one correspondence between processing cells and elements of the computation. For example, an N -body code might associate one processor with each particle. Thus, processors must be allocated on demand, much as memory is allocated in conventional computers. Second, to implement nontrivial algorithms, processors must be able to communicate with each other in patterns which may depend on both the basic algorithm and the dynamic state of an ongoing calculation. For example, when simulating an electronic circuit, processors would communicate with each other in a fixed pattern isomorphic to the simulated circuit diagram; the parallel tree code described in § 4 requires processors to communicate with each other in a constantly changing pattern as the calculation proceeds. Programmable communications patterns are essential for a general-purpose parallel machine.

The Connection Machine * is a requirements in mind. The hardware simple arithmetic-logic unit and 4,096 instruction sequence broadcast from Most instructions, however, are *conditional*: certain flag is true. This flag may be to the result of a local computation or packet-switching communications ne other. For a detailed description of t

4. A parallel tree-code

The economic disadvantages of circumvented by implementing the f the serial algorithm described by B. construction of an "oct" tree which i The tree may be specified by stipula and each cell is recursively divided i Each cell corresponds to an interior cells of the corresponding cell. An e each particle or interior node within computes integerized coordinates (X bits of these coordinates are shuffled

$$(X_{L-1} \dots X_1 X_0, Y_{L-1} \dots Y_1 Y_0)$$

The relative position of two particles values; specifically, two particles lie of their key values are identical. Pa value. The tree is then constructed fi down to 0. At each iteration, meml left and right neighbors; each group the list, and replaced by a single co The effect is to shrink the initial list When the iteration is finished, the li

Figure 1. The 2-
a unit cell (top)

* Connection Machine is a register

tly distant that the error incurred in serial processor is discussed by Barnes calculated to $\sim 1\%$ accuracy in a tenth CPU time scales very closely to the advantage of a tree-code, however, es of N . For $N \simeq 10^5$ to 10^6 , two processors do not have the floating- in reasonable time. Typical vector or ice numeric operations are interwoven are made larger and larger, less and more is used for memory. In the log N , since the expense of "renting"

The Connection Machine * is a general-purpose parallel computer designed with these requirements in mind. The hardware consists of 65,536 basic processing cells, each of which has a simple arithmetic-logic unit and 4,096 bits of local memory. Processors are controlled by a single instruction sequence broadcast from a program running on a conventional "front-end" computer. Most instructions, however, are *conditional*; they only take effect in those processors where a certain flag is true. This flag may be set or cleared independently in each processor, according to the result of a local computation. Processors are connected together by a high-bandwidth packet-switching communications network which allows any processor to send a message to any other. For a detailed description of the Connection Machine, see Hillis (1985).

4. A parallel tree-code

The economic disadvantages of tree-codes mentioned at the end of § 2 may be effectively circumvented by implementing the force-calculation on a fine-grain parallel architecture. As in the serial algorithm described by Barnes & Hut, the first phase of the force calculation is the construction of an "oct" tree which partitions the system into a nested hierarchy of cubical cells. The tree may be specified by stipulating that the largest "root" cell contains the entire system, and each cell is recursively divided into 8 sub-cells until no cell contains more than one particle. Each cell corresponds to an interior node of the tree; the descendants of a node are the sub-cells of the corresponding cell. An example is shown in fig. 1. A single processor is assigned to each particle or interior node within the tree. To construct the tree, each particle processor first computes integerized coordinates (X, Y, Z) scaled between 0 and $2^L - 1$ within the root cell. The bits of these coordinates are shuffled together to produce a $3L$ -bit key:

$$(X_{L-1} \dots X_1 X_0, Y_{L-1} \dots Y_1 Y_0, Z_{L-1} \dots Z_1 Z_0) \mapsto Z_{L-1} Y_{L-1} X_{L-1} \dots Z_1 Y_1 X_1 Z_0 Y_0 X_0.$$

The relative position of two particles in the final tree may be determined by comparing their key values; specifically, two particles lie within the same cell ℓ levels from the root if the top 3ℓ bits of their key values are identical. Particles are linked into a linear list ordered by increasing key value. The tree is then constructed from the bottom (*i.e.*, leaves) up by iterating for ℓ from $L - 1$ down to 0. At each iteration, members of the ordered list compare their key values with their left and right neighbors; each group of nodes which coincide to level ℓ is identified, linked out of the list, and replaced by a single covering cell, which has the coincident nodes as descendants. The effect is to shrink the initial list of particles into an ever-smaller list of ever-larger sub-trees. When the iteration is finished, the list contains a single node, which is the root cell of the tree.

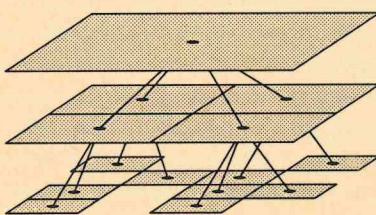


Figure 1. The 2-D analog of an "oct" tree, showing how a unit cell (top) is divided into smaller cells

* Connection Machine is a registered trademark of Thinking Machines Corporation.

Having constructed the tree, the algorithm then computes forces. In a serial tree-code, this is done by recursively descending the tree N times. On a parallel machine, however, it is not acceptable for all N particle processors to simultaneously descend the tree, since this leads to communications bottlenecks near the root. Instead, the tree is used as a network for message distribution. Each particle receives $O(\log N)$ messages, consisting of masses and center-of-mass positions, which describe the rest of the system. This scheme is most easily grasped for a uniformly developed tree, *i.e.*, one which has been subdivided to the same level everywhere; the generalization to non-uniform trees is given below. Such a tree may be viewed as a set of coincident uniform grids, each with spacing twice as fine as the one above it; particles occupy individual cells on the finest level. A particle receives messages from a self-similar hierarchy of ever-coarser grid cells at ever-greater distances, as shown in fig. 2. To insure that the force calculation is correct, individual messages must be received from all adjacent cells on the same level, so that interactions with nearby particles are explicitly included, and the complete hierarchy must cover the entire system without overlap, so that no interactions are missed or counted more than once. Of course, these requirements do not completely specify the hierarchy; the remaining freedom may be used to trade speed for accuracy.

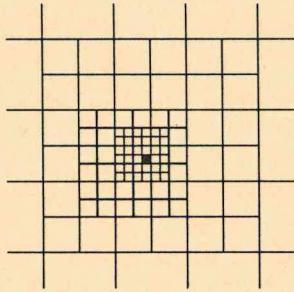


Figure 2. The hierarchy “seen” by a particle in the solid cell

To implement this algorithm efficiently, the tree is augmented with lateral links between each node and its 6 contiguous neighbors on the same level. These links speed message transmission in the lateral direction; more precisely, they are used to setup and advance a “delivery” pointer in each node which points to nearby nodes on the same level. During the force calculation, the delivery pointer traverses a cubical neighborhood of each node. A node which is “geometrically eligible” (see fig. 3) sends its mass and center-of-mass position to the node addressed by the delivery pointer, which then passes the message on to all nodes below it; messages received by particles are used to compute accelerations. The delivery pointers are advanced by short messages sent back from the nodes they currently reference. When the traversal is finished, each particle will have received a series of messages from nodes in the hierarchy above it, providing a concise description of the rest of the system. As a check, each particle computes the total mass of messages received; this should be exactly the total mass of all other particles.

The modifications required to handle non-uniform trees turn out to be relatively simple. Non-uniformity means that some nodes do not exist, and some exist but are isolated, lacking adjacent nodes. Nonexistent nodes neither send nor receive messages; isolated nodes are more problematic, since they must still communicate with other nearby nodes beyond their nonexistent neighbors. When delivery pointers traverse the neighborhood, they must pass over nonexistent

Figure 3. The delivery pointer of the nodes (shaded)

nodes. It is always true that some node is covering a cell, so that instead of accessing this covering cell instead, a delivery pointer can be used to find what its address is. Ideally, these memory pointers would be hardware, which provides the illusion of a much faster memory access than virtual memory provides the illusion of a much faster memory access.

A preliminary version of this code was implemented on the Connection Machine. The machine time requirement is $O(N \log N)$, although for technical reasons it is slower than a parallel direct-summation code for systems with a few thousand particles or more. With a galactic-collision simulation with $N = 10^9$

5. Discussion

This report has shown how, by using a new kind of hardware, large N -body problems can be solved much faster than previously necessary. Specifically, by using $O(N)$ processors, it is possible to solve problems with computing time proportional to $\sim \log N$. While it seems unlikely that a constant-time algorithm will ever be found, there is probably considerable progress to be made.

The algorithm presented here can be applied to a wide range of problems involving interactions and highly inhomogeneous systems. It may be equally well applied to problems in optics, radiative transport, and other areas. The implementation of the algorithm increases the scope and accuracy of the code.

A further implication of this approach is that the physical model is directly reflected in the code. This has been considered an important property in translating fundamentally parallel algorithms into manifest parallelism of the model. The code is therefore harder to debug. This

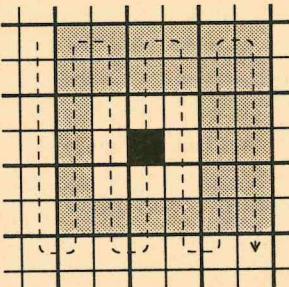


Figure 3. The neighborhood traversed by the delivery pointer of the solid cell. Only geometrically eligible nodes (shaded) are sent messages

nodes. It is always true that some larger cell, which covers the missing node, does exist. By accessing this covering cell instead, a node can determine if the next delivery site exists, and if so, what its address is. Ideally, these modifications are hidden in the kernel of the message-passing code, which provides the illusion of a "virtual tree" with uniform and arbitrarily fine development, much as virtual memory provides the illusion of arbitrarily large address space.

A preliminary version of this force-calculation algorithm has been tested on a Connection Machine. The machine time required to calculate the forces on N particles scales roughly as $O(\log N)$, although for technical reasons this has not yet been tested for $N > 4,096$. Comparison with a parallel direct-summation code shows that the tree algorithm is superior for systems of a few thousand particles or more. With relatively little further work, it should be possible to run a galactic-collision simulation with $N \simeq 6.5 \times 10^4$ in a few hours to a day of machine time.

5. Discussion

This report has shown how, by combining a new kind of force-calculation algorithm with a new kind of hardware, large N -body calculations may be run with far less expenditure than was previously necessary. Specifically, by implementing an $O(N \log N)$ algorithm on a system with $O(N)$ processors, it is possible to solve large- N problems with hardware cost proportional to N in computing time proportional to $\sim \log N$. These scaling laws are probably optimal; for example, it seems unlikely that a constant-time algorithm using $O(N)$ processors exists. On the other hand, there is probably considerable progress to be made in reducing the constants of proportionality.

The algorithm presented here exemplifies a general approach to problems involving long-range interactions and highly inhomogeneous systems. The technique of lumping distant things together may be equally well applied to other kinds of interactions, including electromagnetism and optically-thin radiative transport. The recursive refinement of regions which require more resolution is likewise a general technique with applications in fluid dynamics, image processing, and other areas. The implementation of these techniques on powerful parallel computers significantly increases the scope and accuracy of numerical simulation of realistic physical systems.

A further implication of this approach lies in the fact that the fundamental parallelism of the physical model is directly reflected in the computational implementation. This has not hitherto been considered an important property of computer codes; we have nearly 40 years of experience in translating fundamentally parallel models into algorithms for serial machines. Unfortunately, the manifest parallelism of the model tends to be hidden by the serial implementation, which is therefore harder to debug. This is somewhat like expressing a Lorentz-invariant theory in

non-invariant terms: the symmetry is still there, but may be almost impossible to discern from the formulation. Ultimately, one may want to make the parallelism of computational models as clearly manifest in their implementation as Lorentz invariance is made manifest in the formulation of physical theories. Just as explicit Lorentz invariance is a valuable tool in creating meaningful theories, manifest parallelism may be useful in writing correct algorithms.

I thank Carl Feynman, Daniel Hillis, Piet Hut, and Gerald Sussman for helpful conversations, Thinking Machines Inc. for providing Connection Machine time, and Eric Roberts, DEC-SRC for hospitality during the writing of this report. Support for this research came from NSF grant PHY-8440263.

A GRID

Harvard-Sr

REFERENCES

- Barnes, J. and Hut, P. 1986. "A Hierarchical $O(N \log N)$ Force Calculation Algorithm" *Nature*, in press.
- Hillis, D. 1985. "The Connection Machine" (MIT Press).

The purpose of this paper is to tool, the Ewald method for calculat of Newtonian gravitating particles (this technique to Newtonian cosmic already been made by Efstathiou e potential solver. The present applic high resolution and high accuracy as accomplished with timings proporc unfavorable compared to the $N \log N$ argue that special properties of the it an attractive competitor to P^3M

Suppose we have a system of N cubic lattice. Let the fundamental c cube be represented as

The periodic extension of this funct

where κ runs over all vectors with in

$$\bar{\rho}_\kappa$$

In terms of these coefficients, the po

$$\phi(\mathbf{r})$$

where the Fourier transform of the $\overline{G}(\mathbf{k}) = 4\pi/k^2$. The prime on the su from the cancellation of the mean e.g., Peebles 1980; pp. 41–45).

In Ewald's method the potentia ration and Fourier spaces. This is a and short range parts,

A Parallel Hashed Oct-Tree N-Body Algorithm

Best Student Paper, Supercomputing '93.

Michael S. Warren*
Theoretical Astrophysics
Mail Stop B288
Los Alamos National Laboratory
Los Alamos, NM 87545

John K. Salmon
Physics Department
206-49
California Institute of Technology
Pasadena, CA 91125

Abstract

We report on an efficient adaptive N-body method which we have recently designed and implemented. The algorithm computes the forces on an arbitrary distribution of bodies in a time which scales as $N \log N$ with the particle number. The accuracy of the force calculations is analytically bounded, and can be adjusted via a user defined parameter between a few percent relative accuracy, down to machine arithmetic accuracy. Instead of using pointers to indicate the topology of the tree, we identify each possible cell with a key. The mapping of keys into memory locations is achieved via a hash table. This allows the program to access data in an efficient manner across multiple processors. Performance of the parallel program is measured on the 512 processor Intel Touchstone Delta system. We also comment on a number of wide-ranging applications which can benefit from application of this type of algorithm.

1 Introduction

N-body simulations have become a fundamental tool in the study of complex physical systems. Starting from a basic physical interaction (e.g., gravitational, Coulombic, Biot-Savart, van der Waals) one can follow the dynamical evolution of a system of N bodies, which represent the phase-space density distribution of the system. N-body simulations are essentially statistical in nature (unless the physical system can be directly modeled by N bodies, as is the case in some molecular dynamics simulations). More bodies implies a more accurate and complete sampling of the phase space, and hence more accurate or complete results. Unfortunately, the minimum accuracy required to model systems of interest often depends on having N be much larger than current computational resources allow.

Because interactions occur between each pair of particles in a N-body simulation, the computational work scales asymptotically as N^2 . Much effort has been expended to reduce the computational complexity of such simulations, while retaining acceptable accuracy. One approach is to interpolate the field from a lattice with resolution h , where it can be computed in time $O(h^{-3})$ (using multigrid) or $O(h^{-3} \log h^{-3})$ (using Fourier transforms). The N -dependence of the time complexity then becomes $O(N)$. The drawback to this method is that dynamics on scales comparable to or smaller than h cannot be modeled. In three dimensions, this restricts the dynamic range in length to about one part in a hundred (or perhaps one part in a thousand on a parallel supercomputer), which is insufficient for many calculations.

Another approach is to divide the the interactions into “near” and “far” sets. The forces due to distant particles can then be updated less frequently, or their forces can be ignored completely if one decides the effects of distant particles are negligible. However, this risks significant errors which are hard to analyze. Also, any significant clustering in the system will reduce the efficiency of the method, since a large fraction of the particles end up being in a “near” set. Over the past several years, a number of methods have been introduced which allow N-body simulations to be performed on arbitrary collections of bodies in time much less than $O(N^2)$, without imposition of a lattice. They all have in common the use of a truncated expansion (e.g., Taylor expansion, Legendre expansion, Poisson expansion) to *approximate* the contribution of many bodies with a single *interaction*. The resulting complexity is usually cited as $O(N)$ or $O(N \log N)$, but a careful analysis of what dependent variables should be held constant (e.g., constant per-timestep error, constant integrated error, constant memory, constant relative error with respect to discreteness noise) often leads to different conclusions about the scaling. In any event, the scaling is a tremendous improvement over $O(N^2)$ and the methods allow accurate

*Department of Physics, University of California, Santa Barbara

computations with vastly larger N .

The basic idea of an N-body algorithm based on a truncated series approximation is to partition an arbitrary collection of bodies in such a manner that the series approximation can be applied to the pieces, while maintaining sufficient accuracy in the force (or other quantity of interest) on each particle. In general, the methods represent a system of N *bodies*¹ in a hierarchical manner by the use of a spatial *tree* data structure. Aggregations of bodies at various levels of detail form the internal nodes of the tree, and are called *cells*. Generally, the expansions have a limited domain of convergence, and even where the infinite expansion converges, the truncated expansion introduces errors of some magnitude. Making a good choice of which cells to interact with, and which to reject as being too inaccurate is critical to the success of these algorithms. The decision is controlled by a function which we shall call the multipole acceptance criterion (MAC). Some of the multipole methods which have been described in the literature are briefly reviewed in the next section.

2 Background

2.1 Multipole Methods

Appel was the first to introduce a multipole method [1]. Appel's method uses a binary tree data structure whose leaves are bodies, and internal nodes represent roughly spherical cells. Some care is taken to construct a "good" set of cells which minimize the higher order multipole moments of the cells. The MAC is based on the size of interacting cells. The method was originally thought to be $O(N \log N)$, but has more recently been shown to be $O(N)$ [2].

The Barnes-Hut (BH) algorithm [3] uses a regular, hierarchical cubical subdivision of space (an oct-tree in three dimensions). A two-dimensional illustration of such a tree (a quad-tree) is show in Fig. 1. Construction of BH trees is much faster than construction of Appel trees. In the BH algorithm, the MAC is controlled by a parameter θ , which requires that the cell size, s , divided by the distance from a particle to the cell center-of-mass be less than θ (which is usually in the range of 0.6–1.0). Cell-cell interactions are not computed, and the method scales as $N \log N$.

The fast multipole method (FMM) of Greengard & Rokhlin [4] has achieved the greatest popularity in the broader population of applied mathematicians and computational scientists. It uses high order multipole expansion

¹We refer to both bodies and particles, which should both be understood to be general "atomic" objects which may refer to a mass element, charge, vortex element, panel, or other quantity subject to a multipole approximation.

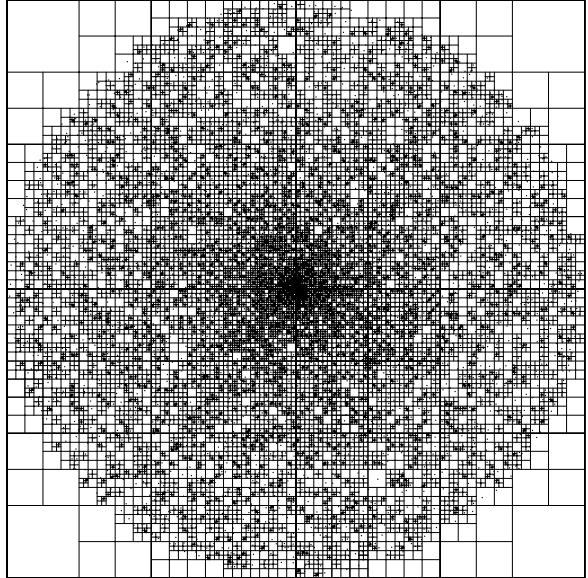


Figure 1: A representation of a regular tree structure in two dimensions (a quad-tree) which contains 10 thousand particles which are centrally clustered.

sions and interacts fixed sets of cells which fulfill the criterion of being "well-separated." The FMM has a well-defined worst case error bound, ϵ , which is guaranteed to be met when multipole expansions are carried out to order $p = -\log_2(\epsilon)$. In two dimensions, when used on systems which are not excessively clustered, the FMM is very efficient. It has been implemented on parallel computers [5, 6]. The crossover point (the value of N at which the algorithm becomes faster than a direct N^2 method) with a stringent accuracy is as low as a few hundred particles. On the other hand, implementations of the FMM in three dimensions have not performed as well. Schmidt and Lee have implemented the algorithm in three dimensions, and find a crossover point of about 70 thousand particles [7]. The reason is that the work in the most computationally intensive step scales as p^2 in two dimensions, and p^4 in three dimensions. It is possible to obtain much better performance by using a smaller p [8, 9], but the worst-case error can become uncomfortably large in this case. The major advantage of the FMM over the methods such as that of Barnes & Hut is that the error bound is rigorously defined. However, this deficiency has been remedied, as is shown in the following section.

2.2 Analytic Error Bounds

Recently, we have analyzed the performance of the Barnes-Hut algorithm, and have shown that the worst case errors can be quite large (in fact, unbounded) for commonly

used values of the opening criterion, θ [10]. We have developed a different method for deciding which cells to interact with. By using moments of the mass or charge distribution within each cell, the method achieves far better worst case error behavior, and somewhat better mean error behavior, for the same amount of computational resources.

In addition, the analysis provides a strict error bound which can be applied to any fast multipole method. This error bound is superior to those used previously because it makes use of information about the bodies contained within a cell. This information takes the form of easily computed moments of the mass or charge distribution (strength) within the cell. Computing this information takes place in the tree construction stage, and takes very little time compared with the later phases of the algorithm. The exact form of the error bound is:

$$\Delta a_{(p)}(\vec{r}) \leq \frac{1}{d^2} \frac{1}{(1 - \frac{b_{max}}{d})^2} \left((p+2) \left(\frac{B_{(p+1)}}{d^{p+1}} \right) - (p+1) \left(\frac{B_{(p+2)}}{d^{p+2}} \right) \right). \quad (1)$$

The moments, $B_{(n)}$ are defined as:

$$B_{(n)} = \int_V d^3x |\rho(x)| |\vec{x} - \vec{r}_0|^n = \sum_{\beta} |m_{\beta}| |\vec{x}_{\beta} - \vec{r}_0|^n. \quad (2)$$

The scalar $d = |\vec{r} - \vec{r}_0|$ is the distance from the particle position \vec{r} to the center of the multipole expansion, p is the largest term in the multipole expansion, and b_{max} is the maximal distance of particles from the center of the cell, (see Fig. 2). This equation is essentially a precise statement of several common-sense ideas. Interactions are more accurate when:

- The interaction distance is larger (larger d).
- The cell is smaller (smaller b_{max}).
- More terms in the multipole expansion are used (larger p).
- The truncated multipole moments are smaller (smaller $B_{(p+1)}$).

Having a per-interaction error bound is an overwhelming advantage when compared to existing multipole acceptance criteria, which assume a worst-case arrangement of bodies within a cell when bounding the interaction error. The reason is that the worst-case interaction error of an *arbitrary* strength distribution is usually many times larger than the error bound on a *particular* strength distribution. This causes an algorithm which knows nothing about the strength distribution inside a cell to provide *too much* accuracy for most multipole interactions. This accuracy is

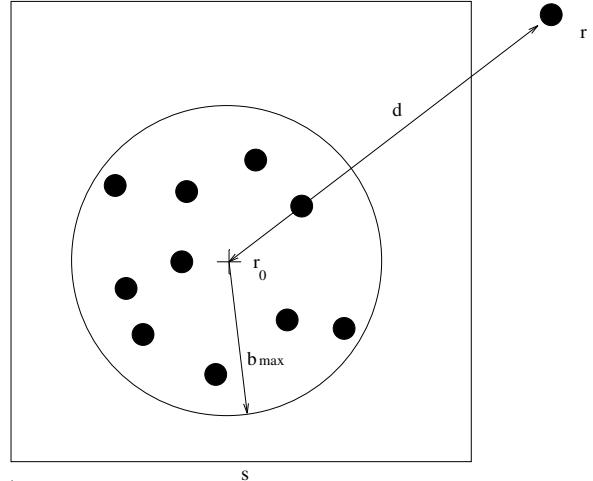


Figure 2: An illustration of the relevant distances used in the error bound equation.

wasted, however, because of the few multipole interaction errors which do approach the worst-case error bound that are added into and pollute the final result. A data-dependent per-interaction error bound is much less prone to this problem, since the resulting error *bound* is much tighter, even though the actual error in the computation is exactly the same.

The implementation of an algorithm using a fixed per-interaction error bound poses little difficulty. One may simply solve for r_c in,

$$\Delta a_{(p)}(r_c) \leq \Delta_{interaction}, \quad (3)$$

where $\Delta_{interaction}$ is a user-specified absolute error tolerance. Then, r_c defines the smallest interaction distance allowed for each cell in the system. For the case of $p = 1$, the critical radius can be analytically derived from Eq. 1 if we use the fact that $B_3 \geq 0$:

$$r_c \geq \frac{b_{max}}{2} + \sqrt{\frac{b_{max}^2}{4} + \sqrt{\frac{3B_2}{\Delta_{interaction}}}}. \quad (4)$$

B_2 is simply the trace of the quadrupole moment tensor. In more general cases (using a better bound on B_3 , or with $p > 1$), r_c can be computed from the error bound equation (Eq. 1) using Newton's method. The overall computational expense of calculating r_c is small, since it need only be calculated once for each cell. Furthermore, Newton's method need not be iterated to high accuracy. The MAC then becomes $d > r_c$ for each displacement d and critical radius r_c (Fig. 3). This is computationally very similar to the Barnes-Hut opening criterion, where instead of using a fixed box size, s , we use the distance r_c , derived from the contents of the cell and the error tolerance. Thus,

our data dependent MAC may replace the MAC in existing algorithms with minimal additional coding.

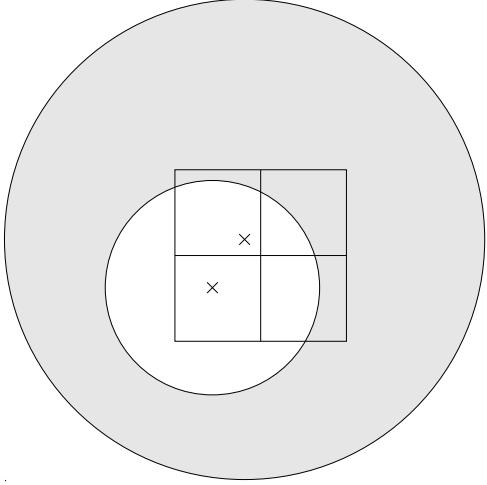


Figure 3: The critical radii of a cell and one of its daughters are shown here as circles. For the specified accuracy, a particle must lie outside the critical radius of a cell. The shaded region shows the spatial domain of all particles which would interact with the lower left daughter cell. Those particles outside the shaded region would interact with the parent cell, and those within the unshaded region would interact with even smaller cells inside the daughter cell.

3 Computational Approach

Parallel treecodes for distributed memory machines are discussed in [11, 12], and their application to the analysis of galaxy formation may be found in [13, 14]. Further analysis and extensions of the computational methods may be found in [15, 16, 17]. The MAC described above is problematical for these previous methods because the parallel algorithm requires determination of locally essential data before the tree traversal begins. With the data-dependent MAC it is difficult to pre-determine which non-local cells are required in advance of the traversal stage. The problem becomes particularly acute if one wishes to impose error tolerances which vary from particle to particle.

It is for this reason that the algorithm described here was developed. It does not rely on the ability to identify *a priori* locally essential data; instead it provides a mechanism to retrieve non-local data as it is needed during the tree traversal. The decision to abandon our previous parallel N-body algorithm was also motivated by the desire to produce a more “friendly” code, with which a variety of research could be performed in computational science as well

as physics. The old code, which was the result of porting a previously existing sequential algorithm, was a maze of complications, brought about by the haphazard addition of pieces over several years. We took full advantage of the opportunity to start over with a clean slate, with the additional benefit of several years of hindsight and experience.

When one considers what additional operations are necessary when dealing with a tree structure distributed over many processors, it is clear that retrieval of particular cells required by one processor from another is a very common operation. When using a conventional tree structure, the pointers in a parent cell in one processor must be somehow translated into a valid reference to daughter cells in another processor. This required translation led us to the conclusion that pointers are not the proper way to represent a distributed tree data structure (at least without significant hardware and operating system support for such operations).

Instead of using pointers to describe the topology of a tree, we use keys and a hash table. We begin by identifying each possible cell with a *key*. By performing simple bit arithmetic on a key, we are able to produce the keys of daughter or parent cells. The tree topology is represented implicitly in the mapping of the cell spatial locations and levels into the keys. The translation of keys into memory locations where cell data is stored is achieved via hash table lookup. Thus, given a key, the corresponding data can be rapidly retrieved. This scheme also provides a uniform addressing mechanism to retrieve data which is in another processor. This is the basis of the hashed oct-tree (HOT) method.

3.1 Key construction and the Hashing Function

We define a key as the result of a map of d floating point numbers (body coordinates in d -dimensional space) into a single set of bits (which is most conveniently represented as a vector of integers). The mapping function consists of translating the floating point numbers into integers, and then interleaving the bits of the d integers into a single key (Fig. 4). Note that we place no restriction on the dimension of the space, although we are physically motivated to pay particular attention to the case of $d = 3$. In this case, the key derived from 3 single precision floating point numbers fits nicely into a single 64 bit integer or a pair of 32 bit integers.

Apart from the trivial choice of origin and coordinate system, this is identical to Morton ordering (also called Z or N ordering, see Chapter 1 of [18] and references therein, and also [19]). This function maps each body in the system to a unique key. We also wish to represent nodes of the tree using this same type of key. In order to distinguish the higher level internal nodes of the tree from the lowest

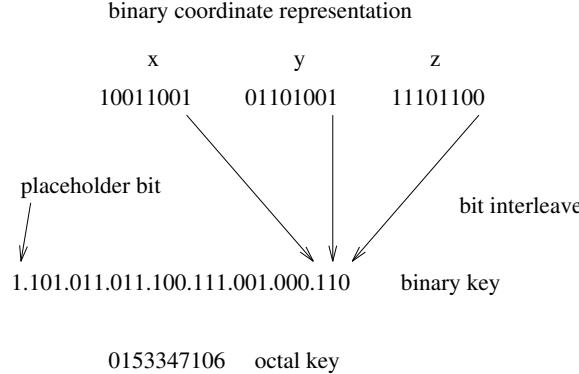


Figure 4: An illustration of the key mapping. Bits of the coordinates are interleaved and a place-holder bit is prepended to the most significant bit. In this example, the 8-bit x , y and z values are mapped to a 25-bit key.

level body nodes, we prepend an additional 1-bit to the most significant bit of every key (the place-holder bit). We may then represent all higher level nodes in the tree in the same key space. Without the place-holder bit, there would be an ambiguity amongst keys whose most significant bits are all zeroes. The root node is represented by the key 1. A two-dimensional representation of such a tree is shown in Fig. 5.

In general, each key corresponds to some composite data describing the physical data inside the domain of a cell (the mass and center-of-mass coordinates, for example). To map the key to the memory location holding this data, a hash table is used. A table with a length much smaller than the possible number of keys is used, with a hashing function to map the k -bit key to the h -bit long hash address. We use a very simple hashing function, which is to AND the key with the bit-mask $2^h - 1$, which selects the least significant h bits.

Collisions in the hash table are resolved via a linked list (chaining). The incidence of collisions could degrade performance a great deal. Our hashing scheme uses the simplest possible function; a one instruction AND. However, it is really the map of floating point coordinates into the key that performs what one usually would consider "hashing." The structure of the hierarchical key space and selection of the least significant bits of the key performs extraordinarily well in reducing the incidence of collisions. For the set of all keys which contain fewer significant bits than the hash mask, the hashing function is "perfect." This set of keys represents the upper levels of the tree, which tend to be accessed the most often. At lower levels of the tree (where the number of bits in a key exceeds the length of the hash mask), distinct keys can result in the same hash address (a collision). However, the map of coordinates

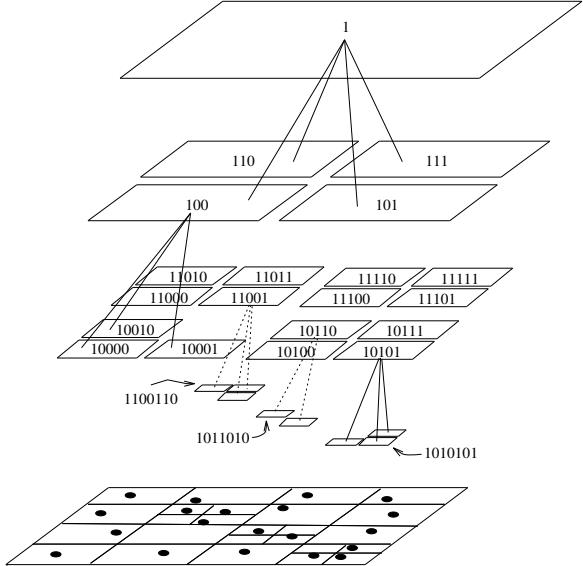


Figure 5: A quad-tree shown along with the binary key coordinates of the nodes. At the bottom is a “flat” representation of the tree topology induced by the 20 particles. The rest of the figure demonstrates the relation of the key coordinates at each level to the tree topology. Many of the links from parent to daughter cells are omitted for clarity.

into the keys keeps these keys spatially separated. On a parallel machine, many of the keys which would result in collisions become distributed to different processors.

The key space is very convenient for tree traversals. In order to find daughter nodes, the parent key is left-shifted by d bits, and the result is added (or equivalently OR'ed) to daughter numbers from 0 to $2^d - 1$. Also, the key retrieval mechanism is much more flexible in terms of the kinds of accesses which are allowed. If we wish to find a particular node of a tree in which pointers are used to traverse the tree, we must start at the root of the tree, and traverse until we find the desired node (which takes of order $\log N$ operations). On the other hand, a key provides immediate ($O(1)$) access to any object in the tree.

An entry in the hash table (an `hcell`) consists of a pointer to the cell or body data, a pointer to a linked list which resolves collisions, the key, and various flags which describe properties of the `hcell` and its corresponding cell. In order to optimize certain tree traversal operations, we also store in each `hcell` 2^d bits that describe which daughters of the cell actually exist. This redundant information allows us to avoid using hash-table lookup functions to search for cells which don't exist.

The use of a hash table offers several important advantages. First, the access to data takes place in a manner which is easily generalized to a global accessing scheme

implementable on a message passing architecture. That is, non-local data may be accessed by requesting a key, which is a uniform addressing scheme, regardless of which processor the data is contained within. This type of addressing is not possible with normal pointers on a distributed memory machine. We can also use the hash table to implement various mechanism for caching non-local data and improving memory system performance.

3.2 Tree Construction

The higher level nodes in the tree can be constructed in a variety of ways. The simplest is analogous to that which was described in [3]. Each particle is loaded into the tree by starting at the root, and traversing the partially constructed tree. When two particles fall within the same leaf node, the leaf is converted to a cell which is entered into the hash table, and new leaves are constructed one level deeper in the tree to hold each of the particles. This takes $O(\log N)$ steps per particle insertion. After the topology of the tree has been constructed, the contents (mass, charge, moments, etc.) of each cell may be initialized by a post-order tree traversal.

A faster method is possible by taking advantage of the spatial ordering implied in the key map. We first sort the body keys, and then consider the bodies in this list in order. As bodies are inserted into the tree, we start the traversal at the location of the last node created (rather than at the root). With this scheme, the average body insertion requires $O(1)$ time. We still require $O(N \log N)$ time to sort the list in the first place, but keeping the body list sorted will facilitate our parallel data decomposition as well. Furthermore, it is also usually the case that numerous tree constructions will take place in the course of a simulation. Since the positions of the bodies in the key space typically do not change a great deal between timesteps, one can take advantage of a more efficient $O(N)$ insertion sort [20], of the almost-sorted data after the first timestep.

3.3 Parallel Data Decomposition

The parallel data decomposition is critical to the performance of a parallel algorithm. A method which may be conceptually simple and easy to program may result in load imbalance which is unacceptable. A method which attempts to balance the work precisely may take so long that performance of the overall application suffers.

We have implemented a method which can rapidly domain decompose a d -dimensional set of particles into load balanced spatial groups which represent the domain of each processor. We take advantage of the properties of the mapping of spatial coordinates to keys to produce a “good”

domain decomposition. The idea is to simply cut the one-dimensional list of sorted body key ordinates (see Fig. 6) into N_p (number of processors) equal pieces, weighted by the amount of work corresponding to each body. The work for each body is readily approximated by counting the number of interactions the body was involved in on the previous timestep. This results in a spatially adaptive decomposition, which gives each processor an equal amount of work. Additionally, the method keeps particles spatially grouped, which is very important for the efficiency of the traversal stage of the algorithm, since the amount of non-local data needed is roughly proportional to the surface area of the processor domain. An illustration of this method on a two-dimensional set of particles is illustrated in Fig. 8 for a highly clustered set of particles (that which was shown in Fig. 1) with $N_p = 16$. One source of inefficiency in the Morton ordered decomposition is that a processor domain can span one of the spatial discontinuities. A possible solution is to use Peano-Hilbert ordering for the domain decomposition, which does not contain spatial discontinuities, which is shown in Fig. 7.

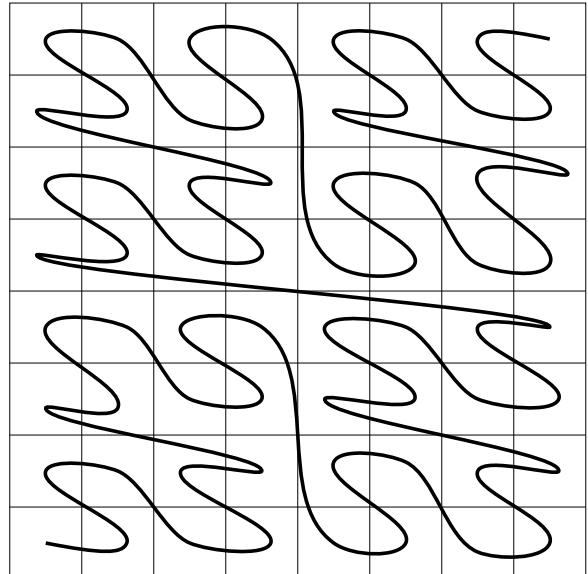


Figure 6: The path indicates the one-dimensional symmetric self-similar path which is induced by the map of interleaved bits (Morton order). The domain decomposition is achieved by cutting the one-dimensional list into N_p pieces.

3.4 Parallel Tree Construction

After the domain decomposition, each processor has a disjoint set of bodies. The initial stage in parallel tree building is the construction of a tree made of the local bodies.

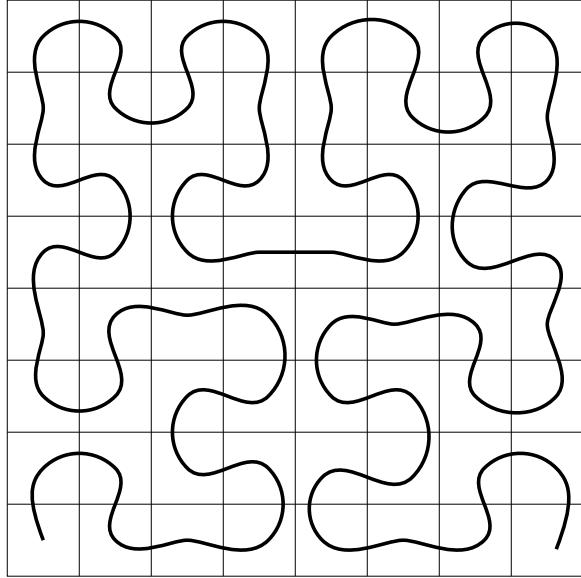


Figure 7: A non-symmetric path that does not contain discontinuities (Peano-Hilbert order) which produces a better decomposition, but in practice does not seem to provide much of a performance increase.

A special case occurs at each processor boundary in the one-dimensional sorted key list, where the terminal bodies from adjacent processors could lie in the same cell. This is taken care of by sending a copy of each boundary body to the adjacent processor, which allows the construction of the proper tree nodes. Then, copies of *branch* nodes from each processor are shared among all processors. This stage is made considerably easier and faster since the domain decomposition is intimately related to the tree topology (unlike the orthogonal recursive bisection method used in our previous code [12]). The branches make up a complete set of cells which represent the entire processor domain at the coarsest level possible. These branch cells are then globally communicated among the processors. All processors can then “fill in” the missing top of the tree down to the branch cells. The address of the processor which owns each branch cell is passed to the destination processor, so the hcell created is marked with its origin. A traversal routine can then immediately determine which processor to request data from when it needs access to the daughters of a branch cell. The daughters received from other processors are also marked in the same fashion. We have also tried implementing the branch communication step in a more computationally clever manner which does not globally concatenate the branches, but its complexity has tended to outweigh its benefit. This does not rule out the possibility of finding a better method for this stage of the algorithm, however.

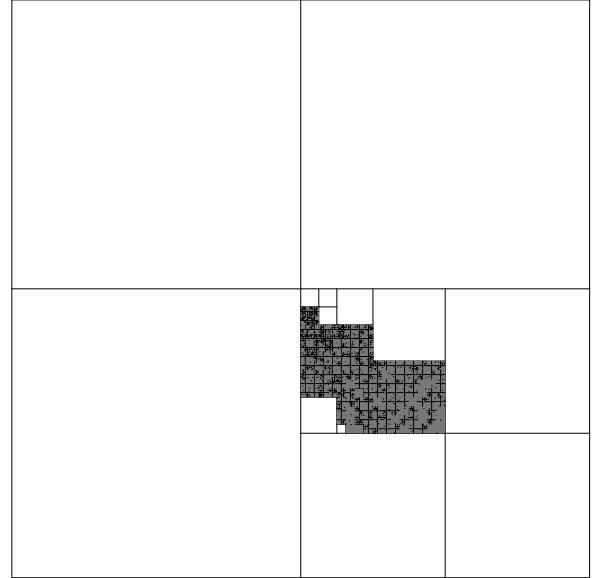


Figure 8: A processor domain for one of 16 processors in a data decomposition for the clustered system of bodies shown in Fig. 1. The domain shown is a result of the decomposition strategy outlined in the text.

3.5 Tree Traversal

A tree traversal routine may be cast in recursive form in a very few lines of C code:

```
Traverse(Key_t key, int (*MAC)(hcell *),
         void (*postf)(hcell *)) {
    hcell *pp;
    unsigned int child;

    if ((pp=Find(key)) && MAC(pp)) return;
    key = KeyLshift(key, NDIM);
    for (child = 0; child < (1<<NDIM); child++)
        Traverse(KeyOrInt(key, child), MAC, postf);
    postf(pp);
}
```

This code applies an arbitrary MAC to determine whether to continue traversing the children of a cell. If the children are traversed, than another function, postf, is called upon completion of the descendants. By appropriate choice of the MAC and postf one can execute pre-order or post-order traversals with or without complex pruning strategies (i.e., multipole acceptability criteria).

On a parallel machine, one may add additional functionality to the Find function, in order to handle cases where the requested node is in the memory of another processor. The additional code would request non-local data, wait to receive it, and insert it into the tree. This allows the same traversal code fragment to work without further modification on a distributed memory computer. However,

the performance of such an approach is bound to be dismal. Each request of non-local data is subject to the full interprocessor communication latency. Computation stalls while waiting for the requested data to arrive.

It is possible to recast the traversal function in a form which allows the entire *context* of the traversal to be stored. In this case, when a request for non-local data is encountered, the request is buffered, and the computation may proceed. Almost all of the latency for non-local data requests may be hidden, by trading communication latency for a smaller amount of complexity overhead.

The traversal method we have chosen is breadth-first list based scheme. It does not use recursion, and has several useful properties. We shall discuss the plain sequential method first, and then show the additions to allow efficient traversals on a parallel machine.

The input to the list-based traversal is a *walk list* of hcell nodes. On the first pass, the walk list contains only the root hcell. Each daughter of the input walk list nodes is tested against the MAC. If it passes the MAC, the corresponding cell data is placed on the *interaction list*. If a daughter fails the MAC, it is placed on the output walk list. After the entire input list is processed the output walk list is copied to the walk list and the process iterates. The process terminates when there are no nodes remaining on the walk list. This method has an advantage over a recursive traversal in that there is an opportunity to do some vectorization of the intermediate traversal steps, since there are generally a fair number of nodes which are being tested at a time. It also results in a final interaction list which can be passed to a fully vectorized force calculation routine. The details are too intricate to allow us to present real C code, so we present the algorithm in pseudocode instead:

```
ListTraverse((*MAC)(hcell *))
{
    copy root to walk_list;
    while (!Empty(walk_list)) {
        for (each item on walk_list) {
            for (each daughter of item) {
                if (MAC(daughter))
                    copy daughter to interact_list;
                else
                    copy daughter to output_walk_list;
            }
        }
        walk_list = output_walk_list;
    }
}
```

When the traversal is complete, the *interact_list* contains a vector of items that must undergo interactions (according to the particular MAC). The interactions themselves may be computed separately, so that code may be vectorized and optimized independently of the tree traversal method.

3.6 A Latency Hiding Tree Traversal

On a parallel machine, the traversal will encounter hcells for which the daughters are not present in local memory. In this case we add some additional lists which allow computation to proceed, while the evaluation of the non-local data is deferred to some later time. Each hcell is labeled with a HERE bit. This bit is set if the daughters of the hcell are present in local memory. This bit is tested in the traversal before the attempt to find the daughters. If the HERE bit is not set, the key and the source processor address (which is contained in the hcell) are placed on the *request list*, and another copy of the key is placed on a *defer list*. We additionally set a REQUESTED bit in the hcell, to prevent additional requests for the same data. This allows processing to continue on the hcells in the input walk list. As the traversal proceeds, additional requests will occur, until a final state is reached, where as much progress as possible has been made on the given traversal (using only data in local memory). In this state, there are a number of keys and processor addresses in the request list, and an equal number of keys in the defer list, which require non-local data to be received before the traversal may continue.

The request list is periodically translated into a series of interprocessor messages which contain requests for data. Upon receipt of such a message, the appropriate hcells are packaged into a reply, and the answer is returned via a second interprocessor message. When a reply is received, an appropriate entry is made in the hash table, and subsequent Find requests will return the data. It is possible to implement this request/reply protocol either loosely synchronously or asynchronously. The decision is governed by the level of support and relative performance offered by the hardware and operating system.

Upon receipt of some replies (it is not necessary to wait for all replies to arrive), the defer list can be renamed as the *walk_list*, and the traversal can be restarted with the newly arrived data. Alternatively, one can begin an entirely separate traversal to compute, e.g., the force on another particle. With appropriate bookkeeping one can tolerate very long latencies by implementing a circular queue of active traversals (with a shared request list). We have used a circular queue with 30 active traversals, so that after 30 traversals have been deferred, we restart the first traversal by copying its defer list to its walk list. The requested data has usually arrived in the interim.

3.7 Memory Hierarchy and Access Patterns

Treecodes place heavy demands on the memory subsystems of modern computers. The quasi-random access to widely separated memory locations during the tree traversal receives little benefit from a small on-chip cache, and

can in fact overwhelm the translation look-a-side buffer (TLB) on microprocessors similar to the i860. This results in very poor performance from algorithms which have been designed without consideration of memory bandwidth and the limitations inherent in memory hierarchies which are intended to function with DRAM significantly slower than the processor cycle time.

Since the tree traversal is so stressful to most memory architectures, we have arranged the order of computation to take advantage of the underlying structure of the algorithm, which helps encourage a more orderly memory access pattern. A useful property of treecode algorithms is that particles which are spatially near each other tend to have very similar cell interaction lists. By updating the particles in an order which takes advantage of their spatial proximity, we can reduce the number of memory accesses which miss the cache and TLB. A convenient and efficient ordering once again uses the same sorted key list used in the tree construction. By updating particles in the order defined by the key map (Fig. 6), we achieve this goal.

An additional technique to improve memory access speed is through the rearrangement of data in the linked list of collisions in the hash table. By moving data which has been recently accessed to the top of the linked list, it is possible to create a “virtual cache” by keeping often used data in the contiguous memory locations making up the hash table. This also allows one to obtain good performance with a hash table much smaller than one would naively expect.

The more extended memory hierarchy in a distributed memory parallel computer (possibly with virtual memory on each node) can benefit from this scheme as well. We wish to keep things for as long as possible in the fastest level of the hierarchy that includes registers, cache, local memory, other processors memory, and virtual memory. We could extend the “virtual cache” model even further, by erasing data which has come from another processor which has not been used recently. Although this has not been implemented, we expect that it will allow significantly larger simulations to take place, since the majority of the memory used consists of copies of cells from other processors.

4 Performance

Here we provide timings for the various stages of the algorithm on the 512 processor Intel Touchstone Delta installed at Caltech. The timings listed are from an 8.8 million particle production run simulation involving the formation of structure in a cold dark matter Universe [14]. During the initial stages of the calculation, the particles are spread uniformly throughout the spherical computational volume. We set an absolute error bound on each partial acceleration

of 10^{-3} times the mean acceleration in the system. This results in 2.2×10^{10} interactions per timestep in the initial unclustered system. The timing breakdown is as follows:

<i>computation stage</i>	<i>time (sec)</i>
Domain Decomposition	7
Tree Build	7
Tree Traversal	33
Data Communication During Traversal	6
Force Evaluation	54
Load Imbalance	7
Total (5.8 Gflops)	114

At later stages of the calculation the system becomes extremely clustered (the density in large clusters of particles is typically 10^6 times the mean density). The number of interactions required to maintain the same accuracy grows moderately as the system evolves. At a slightly increased error bound of 4×10^{-3} , the number of interactions in the clustered system is 2.6×10^{10} per timestep.

<i>computation stage</i>	<i>time (sec)</i>
Domain Decomposition	19
Tree Build	10
Tree Traversal	55
Data Communication during traversal	4
Force Evaluation	60
Load Imbalance	12
Total (4.9 Gflops)	160

It is evident that the initial domain decomposition and tree building stages take a relatively larger fraction of the time in this case. The reason is that in order to load balance the force calculation, some processors have nearly three times as many particles as the mean value, and over ten times as many particles as the processor with the fewest. The load balancing scheme currently attempts to load balance only the work involved in force evaluation and tree traversal, so the initial domain decomposition and tree construction work (which scales closely with the particle number within the processor) becomes imbalanced.

Note that roughly 50% of the execution time is spent in the force calculation subroutine. This routine consists of a few tens of lines of code, so it makes sense to obtain the maximum possible performance through careful tuning. For the Delta’s i860 microprocessor we used hand coded assembly language to keep the three-stage pipeline fully filled, which results in a speed of 28 Mflops per processing node in this routine.

If we count only the floating point operations performed in the force calculation routine as “useful work” (30 flops per interaction) the overall speed of the code is about 5–6 Gflops. However, this number is in a sense unfair to

the overall algorithm, since the majority of the code is not involved in floating point operations at all, but with tree traversal and data structure manipulation. The integer arithmetic and addressing speed of the processor are as important as the floating point performance. We hope that in the future, evaluation of processors does not become overbalanced toward better floating point speed at the expense of integer arithmetic and memory bandwidth, as this code is a good example of why a balanced processor architecture is necessary for good overall performance.

5 Multi-purpose Applications

Problems of current interest in a wide variety of areas rely heavily on N-body and/or fast multipole methods. Accelerator beam dynamics, astrophysics (galaxy formation, large-scale structure), computational biology (protein folding), chemistry (molecular structure and thermodynamics), electromagnetic scattering, fluid mechanics (vortex method, panel method), molecular dynamics, and plasma physics, to name those we are familiar with, but there are certainly more. In some of these areas, N^2 algorithms are still the most often used, due to their simplicity. However, as problems grow larger, the use of fast methods becomes a necessity. Indeed, in the case of problems such as electromagnetic scattering, a fast multipole method reduces the operation count for solving the second-kind integral equation from $O(N^3)$ for Gaussian elimination to $O(N^{4/3})$ per conjugate-gradient iteration [21]. Such a vast improvement allows one to contemplate problems which were heretofore simply impossible. Alternatively, one can use a workstation to solve problems that had previously been in the sole domain of large supercomputers.

We have spent substantial effort in this code keeping the data structures and functions required by the “application” away from those of the “tree”. With suitable abstractions and ruthless segregation, we have met with some success in this area. We currently have a number of physics applications which share the same tree code. In general, the addition of another application only requires the definition of a data structure, and additional code is required only with respect to functions which are physics related (e.g., the force calculation).

We have described the application of our code to gravitational N-body problems above. The code has also been indispensable in performing statistical analyses and data processing on the end result of our N-body calculations, since their size prohibits analysis on anything but a parallel supercomputer. The code also has a module which can perform three-dimensional compressible fluid dynamics using smoothed particle hydrodynamics (with or without gravity). We have also implemented a vortex particle method

[22]. It is a simple matter to use the same program to do physics involving other force laws. Apart from the definition of a data structure and modification of the basic force calculation routine, one only need derive the appropriate MAC using the method described in Salmon & Warren [10].

6 Future Improvements

The code described here is by no means a “final” version. The implementation has been explicitly designed to easily allow experimentation, and inclusion of new ideas which we find useful. It is perhaps unique in that it is serving double duty as a high performance production code to study the process of galaxy formation, as well as a testbed to investigate multipole algorithms.

Additions to the underlying method which we expect will improve its performance even further include the addition of cell-cell evaluations (similar to those used in the fast multipole method) and the ability to evolve each particle with an independent timestep (which improves performance significantly in systems where the timescale varies greatly). We expect that the expression of the algorithm in the C++ language will produce a more friendly program by taking advantage of the features of the language such as data abstraction and operator overloading. The code is very portable to other parallel platforms, and we currently have code running on the Intel Paragon, the CM-5, the IBM SP-1, and networks of workstations. The bulk of the remaining improvements are in the area of processor specific tuning, such as CDPEAC coding of the inner loop of the force-evaluation routine to obtain optimal floating point performance on the CM-5.

7 Conclusion

In an overall view of this algorithm, we feel that these general items deserve special attention:

- The fundamental ideas in this algorithm are, for the most part, standard tools of computer science (key mapping, hashing, sorting). We have shown that in combination, they form the basis of a clean and efficient parallel algorithm. This type of algorithm does not evolve from a sequential method. It requires starting anew, without the prejudices inherent in a program (or programmer) accustomed to using a single processor.
- The raw computing speed of the code on an extremely irregular, dynamically changing set of particles which

require global data for their update, using a large number of processors (512), is comparable with the performance quoted for much more regular static problems, which are sometimes identified as the only type of “scalable” algorithms which obtain good performance on parallel machines. We hope we have convinced the reader that even difficult irregular problems are amenable to parallel computation.

We expect that algorithms such as that described here, coupled with the extraordinary increase in computational power expected in the coming years, will play a major part in the process of understanding complex physical systems.

Acknowledgments

We thank Sanjay Ranka for pointing out the utility of Peano-Hilbert ordering. We thank the CSCC and the CCSF for providing computational resources. JS wishes to acknowledge support from the Advanced Computing Division of the NSF, as well as the CRPC. MSW wishes to acknowledge support from IGPP and AFOSR. This research was supported in part by a grant from NASA under the HPCC program. This research was performed in part using the Intel Touchstone Delta System operated by Caltech on behalf of the Concurrent Supercomputing Consortium.

References

- [1] A. W. Appel, “An efficient program for many-body simulation,” *SIAM J. Computing*, vol. 6, p. 85, 1985.
- [2] K. Esselink, “The order of Appel’s algorithm,” *Information Processing Let.*, vol. 41, pp. 141–147, 1992.
- [3] J. Barnes and P. Hut, “A hierarchical O(NlogN) force-calculation algorithm,” *Nature*, vol. 324, p. 446, 1986.
- [4] L. Greengard and V. Rokhlin, “A fast algorithm for particle simulations,” *J. Comp. Phys.*, vol. 73, pp. 325–348, 1987.
- [5] L. Greengard and W. D. Gropp, “A parallel version of the fast multipole method,” *Computers Math. Applic.*, vol. 20, no. 7, pp. 63–71, 1990.
- [6] F. Zhao and S. L. Johnsson, “The parallel multipole method on the connection machine,” *SIAM J. Sci. Stat. Comp.*, vol. 12, pp. 1420–1437, Nov. 1991.
- [7] K. E. Schmidt and M. A. Lee, “Implementing the fast multipole method in three dimensions,” *J. Stat. Phys.*, vol. 63, no. 5/6, pp. 1223–1235, 1991.
- [8] J. A. Board, J. W. Causey, J. F. Leathrum, A. Windemuth, and K. Schulten, “Accelerated molecular dynamics simulation with the parallel fast multipole algorithm,” *Chem. Phys. Let.*, vol. 198, p. 89, 1992.
- [9] H.-Q. Ding, N. Karasawa, and W. Goddard, “Atomic level simulations of a million particles: The cell multipole method for coulomb and london interactions,” *J. of Chemical Physics*, vol. 97, pp. 4309–4315, 1992.
- [10] J. K. Salmon and M. S. Warren, “Skeletons from the treecode closet,” *J. Comp. Phys.*, 1992. (in press).
- [11] J. K. Salmon, *Parallel Hierarchical N-body Methods*. PhD thesis, California Institute of Technology, 1990.
- [12] M. S. Warren and J. K. Salmon, “Astrophysical N-body simulations using hierarchical tree data structures,” in *Supercomputing ’92*, IEEE Comp. Soc., 1992. (1992 Gordon Bell Prize winner).
- [13] M. S. Warren, P. J. Quinn, J. K. Salmon, and W. H. Zurek, “Dark halos formed via dissipationless collapse: I. Shapes and alignment of angular momentum,” *Ap. J.*, vol. 399, pp. 405–425, 1992.
- [14] W. H. Zurek, P. J. Quinn, J. K. Salmon, and M. S. Warren, “Formation of structure in a CDM universe: Correlations in position and velocity,” *Ap. J.*, 1993. (in preparation).
- [15] J. P. Singh, J. L. Hennessy, and A. Gupta, “Implications of hierarchical N-body techniques for multiprocessor architectures,” Tech. Rep. CSL-TR-92-506, Stanford University, 1992.
- [16] J. P. Singh, C. Holt, T. Totsuka, A. Gupta, and J. L. Hennessy, “Load balancing and data locality in hierarchical N-body methods,” *Journal of Parallel and Distributed Computing*, 1992. (in press).
- [17] S. Bhatt, M. Chen, C. Y. Lin, and P. Liu, “Abstractions for parallel N-body simulations,” Tech. Rep. DCS/TR-895, Yale University, 1992.
- [18] H. Samet, *Design and Analysis of Spatial Data Structures*. Reading, MA: Addison-Wesley, 1990.
- [19] J. E. Barnes, “An efficient N-body algorithm for a fine-grain parallel computer,” in *The Use of Supercomputers in Stellar Dynamics* (P. Hut and S. McMillan, eds.), (New York), pp. 175–180, Springer-Verlag, 1986.
- [20] D. E. Knuth, *The Art of Computer Programming: Sorting and Searching*, vol. 3. Reading, Mass.: Addison Wesley, 1973.
- [21] N. Engheta, W. D. Murphy, V. Rokhlin, and M. S. Vassiliou, “The fast multipole method (FMM) for electromagnetic scattering problems,” *IEEE Transactions on Antennas and Propagation*, vol. 40, no. 6, pp. 634–642, 1992.
- [22] J. K. Salmon, M. S. Warren, and G. S. Winckelmans, “Fast parallel tree codes for gravitational and fluid dynamical N-body problems,” *International Journal of Supercomputing Applications*, 1993. (submitted).

nment function when P would be in the reference
rectional errors could lead to total or equivalently

(8-26)

(8-27)

source of errors. interlacing or a small errors in the force beyond the ratio to that of E_{ref} main factors invisible by using a force or larger Q . The practical

PM calculation.

$$\sum_n |\hat{\mathbf{R}}|^2 \quad (8-28)$$

values \mathbf{n}' left after sum over \mathbf{n}' values (28) are the same comes

(8-29)

(8-30)

8-4 THE SHORT-RANGE FORCE

The total short-range part of the force on a particle i at position \mathbf{x}_i is given by the sum of the interparticle short-range forces

$$\mathbf{F}_i^{\text{sr}} = \sum_{j=1}^{N_p} \mathbf{f}_{ij}^{\text{sr}} \quad (8-31)$$

The elementary method of evaluating \mathbf{F}_i^{sr} is to sweep through all particles $j = 1, \dots, N_p$, test whether the separation $r_{ij} = |\mathbf{x}_i - \mathbf{x}_j|$ is less than r_e , and, if so, compute $\mathbf{f}_{ij}^{\text{sr}}$ and add it to \mathbf{F}_i^{sr} . Such an approach is clearly impractical, since for each of the N_p values of i one would have to test $N_p - 1$ separations r_{ij} giving an operations count scaling as N_p^2 .

8-4-1 The Chaining Mesh

The computational cost of locating those particles j which contribute to the short-range force on particle i is greatly reduced if the particle coordinates are ordered such that the tests for locating particles j such that $r_{ij} \leq r_e$ need only be performed over a small subset N_n of the total number of particles N_p . It is for this reason that the chaining mesh is introduced. The chaining mesh (in three dimensions) is a regular lattice of $(M_1 \times M_2 \times M_3)$ cells, covering the computational box (of side $L_1 \times L_2 \times L_3$) in much the same manner as the $(N_1 \times N_2 \times N_3)$ cells of the much

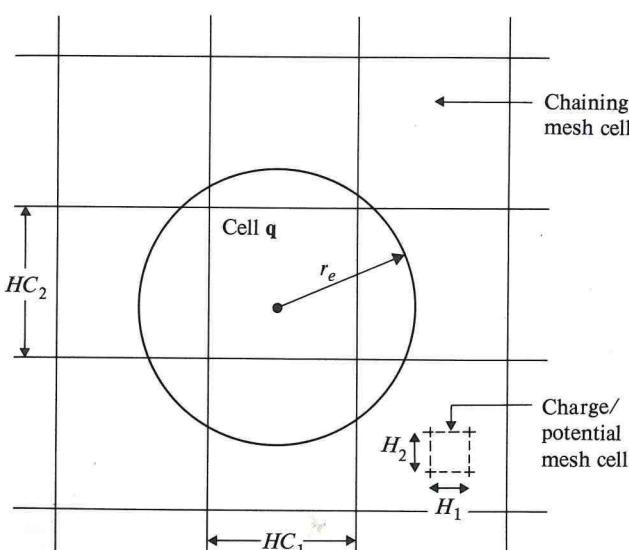


Figure 8-2 In the short-range force calculation, the computational box is divided into chaining cells. Contributions to the force \mathbf{F}_i^{sr} on particle i in cell q are nonzero only from particles j in cell q and the neighboring cells.

finer charge-potential mesh. The number of cells M_s along the s direction is given by the largest integer less than or equal to L_s/r_e . Consequently, the lengths of the sides of the cells of the chaining mesh are always greater than or equal to the cutoff radius r_e .

Figure 8-2 depicts a chaining mesh in two dimensions. Typically, the side lengths of the chaining mesh cells HC_s are between three and four times greater than the side lengths H_s of the cells of the charge-potential mesh. The circle of radius r_e centered on particle i in chaining cell \mathbf{q} delineates the area in which particles j must lie if they are to have a nonzero contribution to \mathbf{F}_i^{sr} . Since $HC_s \geq r_e$ for all s , it follows that those particles j which have nonzero contributions to \mathbf{F}_i^{sr} must either lie in the same cell \mathbf{q} as particle i or in one of the eight neighboring cells. If the particle coordinates are sorted into lists for each chaining cell, then to find the force \mathbf{F}_i^{sr} on particle i involves approximately $9N_C$ tests, where $N_C (= N_p/M_1 M_2)$ is the average number of particles per chaining cell. Therefore, if Newton's third law is used, the total number of tests in finding all the short-range forces is approximately $N_n N_p \approx 4.5 N_C N_p$ as compared with N_p^2 for the elementary approach. Similarly, in three dimensions, sorting coordinates into chaining cells gives the number of tests $N_n N_p \sim 13 N_C N_p$.

8-4-2 The Linked Lists

For serial computers (but not necessarily for vector or array processor machines) it is computationally more efficient to sort the coordinate addresses rather than the coordinates themselves. Address sorting is made possible by introducing the linked-list array LL.

If we let $\text{HOC}(\mathbf{q})$ be the head-of-chain table entry for chaining cell \mathbf{q} , and let $\text{LL}(i)$ be the link coordinate for particle i , then the procedure for sorting coordinates into lists for each chaining cells by means of address sorting is summarized as follows:

1. set $\text{HOC}(\mathbf{q}) = 0$ for all \mathbf{q} .
2. do for all particles i .
 - (a) locate cell containing particle
 $\mathbf{q} := \text{int}(x_1/HC_1, x_2/HC_2, x_3/HC_3)$
 - (b) add particle i to head of list for cell \mathbf{q}
 $\text{LL}(i) := \text{HOC}(\mathbf{q})$
 $\text{HOC}(\mathbf{q}) := i$

In two dimensions, the third components of \mathbf{q} and \mathbf{x} are omitted.

The way the sorting procedure works is illuminated by considering an example. Consider the case where three particles i_1, i_2, i_3 lie in chaining cell \mathbf{q} , where $i_1 < i_2 < i_3$. We represent the coordinates by a three-partition box.

i	\mathbf{X}_i	LL_i
-----	----------------	---------------

rection is given
e lengths of the
ial to the cutoff

ically, the side
r times greater
1. The circle of
area in which
Since $HC_s \geq r_e$
tributions to \mathbf{F}_i^{sr}
ght neighboring
ing cell, then to
c tests, where
ell. Therefore, if
the short-range
the elementary
o chaining cells

essor machines)
sses rather than
introducing the

ig cell \mathbf{q} , and let
ure for sorting
dress sorting is

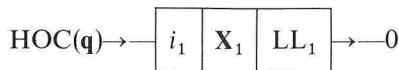
considering an
chaining cell \mathbf{q} ,
on box.

where i is the address (or array element in FORTRAN), \mathbf{X}_i are the physical particle coordinates ($\mathbf{x}_i, \mathbf{P}_i$), and LL_i is the linked-list coordinate. If particle coordinates are swept through in increasing i values then the linked list for cell \mathbf{q} develops as follows:

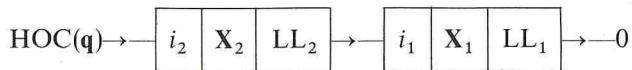
initially (after step 1)

$\text{HOC}(\mathbf{q}) \rightarrow -0$

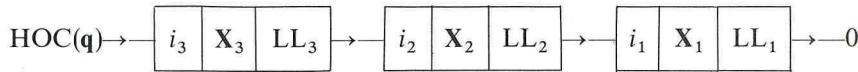
$i_1 < i < i_2$



$i_2 < i < i_3$



$i_3 < i$



The speed and simplicity of creating the linked lists from scratch make it pointless saving and updating them timestep by timestep. The whole sorting process requires only three real arithmetic operations per particle in three dimensions or two in two dimensions.

Once the HOC and LL tables have been filled, a zero entry in $\text{HOC}(\mathbf{q})$ indicates that there are no particles in chaining cell \mathbf{q} . A nonzero entry gives the address of the coordinates of the first particle in the list. The link coordinate of a particle either gives the address of the coordinates of the next particle in the list, or is zero to indicate the end of the list. Therefore, given HOC and LL, coordinates in each cell can be looked up without any searching.

8-4-3 The Momentum Change

The computation of the short-range force contribution to the momentum change at each timestep is a two-stage process. The first stage is to fill the HOC and LL arrays as described above, and the second is to accumulate the changes in momentum.

In three dimensions, particles which have a nonzero contribution to \mathbf{F}_i^{sr} must lie either in the same chaining cell as particle i or in one of the twenty-six neighboring chaining cells. To eliminate the unnecessary double computation of values $\mathbf{f}_{ij}^{\text{sr}}$ used in incrementing the momenta of both particles i and j , we want to process each particle pair (i,j) , $i \neq j$ only once. One method of achieving this is to sweep through each chaining cell, which we will call the current cell. For the current-chaining cell, $\mathbf{f}_{ij}^{\text{sr}}$ is obtained for all pairs (i,j) with both members in the